





The final DRAD-CF-MPC architecture is explained through stepwise defense gains, surrogate fidelity, and AI-learned robustness drivers.



A quantifies architecture gains; B validates surrogate predictability; C ranks learned drivers; D maps algorithmic modules to engineering effects.